

Class 9 | Halloween Mini-Project

AUTHOR

Christopher Leone | A16731724

Background

Today we will take a step back to some data we can taste and explore—we will examine the correlation structure and principal components of some Halloween candy.

```
# Let's read the data into our project:
candy <- read.csv("candy-data.csv", row.names=1)
head(candy)
```

	chocolate	fruity	caramel	peanut	almond	nougat	crisped	rice	wafer
100 Grand	1	0	1		0	0			1
3 Musketeers	1	0	0		0	1			0
One dime	0	0	0		0	0			0
One quarter	0	0	0		0	0			0
Air Heads	0	1	0		0	0			0
Almond Joy	1	0	0		1	0			0

	hard	bar	pluribus	sugar	percent	price	percent	win	percent
100 Grand	0	1	0		0.732	0.860	66.97173		
3 Musketeers	0	1	0		0.604	0.511	67.60294		
One dime	0	0	0		0.011	0.116	32.26109		
One quarter	0	0	0		0.011	0.511	46.11650		
Air Heads	0	0	0		0.906	0.511	52.34146		
Almond Joy	0	1	0		0.465	0.767	50.34755		

(Q1) How many different candy types are in this dataset?

```
# Shows that we have 85 entries.
nrow(candy)
```

```
[1] 85
```

(Q2) How many fruity candy types are in the dataset?

```
# 38 fruity candies!
sum(candy$fruity)
```

```
[1] 38
```

What is your favorite candy?

(Q3) What is your favorite candy in the dataset and what is it's winpercent value?

My favorite Halloween candy are the Nestle Crunch candies. How popular is it?

```
candy["Nestle Crunch",]$winpercent
```

[1] 66.47068

(Q4) What is the winpercent value for "Kit Kat"?

```
candy["Kit Kat",]$winpercent
```

[1] 76.7686

(Q5) What is the winpercent value for "Tootsie Roll Snack Bars"?

```
candy["Tootsie Roll Snack Bars",]$winpercent
```

[1] 49.6535

A helpful function we can use to get a summary of a dataset is the `skim()` function from the **skimr** package, which we downloaded and can call to examine our `candy` dataset.

```
library(skimr)
skim(candy)
```

Data summary	
Name	candy
Number of rows	85
Number of columns	12
<hr/>	
Column type frequency:	
numeric	12
<hr/>	
Group variables	None

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmondy	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

(Q6) Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

The **winpercent** column appears to be on a different scale because it is the only column without values consistently below 1.

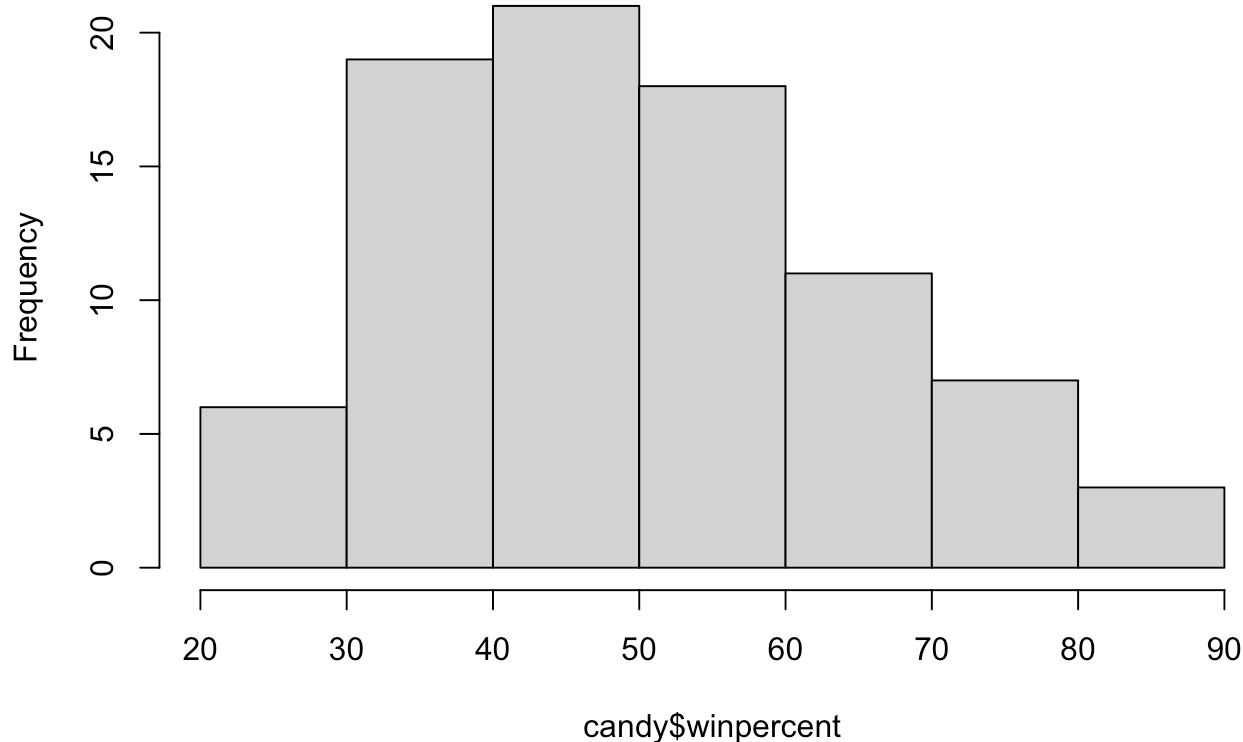
(Q7) What do you think a zero and one represent for the candy\$chocolate column?

Each "zero" represents a "No/FALSE", while each "one" represents a "Yes/TRUE", such as in binary code.

(Q8) Plot a histogram of winpercent values.

```
# Let's just generate a simple histogram here
hist(candy$winpercent)
```

Histogram of candy\$winpercent



(Q9) Is the distribution of winpercent values symmetrical?

No, the distribution is swayed to the left, closer to lower `candy$winpercent` values.

(Q10) Is the center of the distribution above or below 50%?

As alluded to in **Q8**, the center of the distribution is below 50%.

(Q11) On average is chocolate candy higher or lower ranked than fruit candy?

Let's check the values for each:

```
# Here is the mean value for chocolate candy:  
mean(candy$winpercent[as.logical(candy$chocolate)])
```

```
[1] 60.92153
```

```
# And here is the mean for fruity candy:  
mean(candy$winpercent[as.logical(candy$fruity)])
```

```
[1] 44.11974
```

We can see that chocolate candy has ~**16%** higher `winpercent` values than fruity candy.

(Q12) Is this difference statistically significant?

```
# We can use a T-test to find the p-value and significance:  
t.test(x=candy$winpercent[as.logical(candy$chocolate)],y=candy$winpercent[as.logical(cand
```

Welch Two Sample t-test

```
data: candy$winpercent[as.logical(candy$chocolate)] and  
candy$winpercent[as.logical(candy$fruity)]  
t = 6.2582, df = 68.882, p-value = 2.871e-08  
alternative hypothesis: true difference in means is not equal to 0  
95 percent confidence interval:  
 11.44563 22.15795  
sample estimates:  
mean of x mean of y  
 60.92153  44.11974
```

We get a very small p-value of **2.871e-08**, which indicates a very statistically significant difference between the `winpercent` of chocolate and fruity candies.

Overall Candy Rankings

We can make this easier by ordering the data by `winpercent`, which would allow us to see the most/least liked candies.

(Q13) What are the five least liked candy types in this set?

There are two related functions that can help here, one is the classic `sort()` function and the `order()` function. But, here let's play with the `dplyr` package.

```
# Let's experiment with the dplyr package!  
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

`filter`, `lag`

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
candy %>% arrange(winpercent) %>% head(5)
```

	chocolate	fruity	caramel	peanut	almond	nougat
Nik L Nip	0	1	0		0	0
Boston Baked Beans	0	0	0		1	0
Chiclets	0	1	0		0	0
Super Bubble	0	1	0		0	0
Jawbusters	0	1	0		0	0

	crisped	rice	wafer	hard	bar	pluribus	sugar	percent	price	percent
Nik L Nip				0	0	1		0.197		0.976
Boston Baked Beans				0	0	1		0.313		0.511
Chiclets				0	0	1		0.046		0.325
Super Bubble				0	0	0		0.162		0.116
Jawbusters				0	1	0	1	0.093		0.511

	winpercent
Nik L Nip	22.44534
Boston Baked Beans	23.41782
Chiclets	24.52499
Super Bubble	27.30386
Jawbusters	28.12744

The 5 least liked candies in the dataset are "Nik L Nip", "Boston Baked Beans", "Chiclets", "Super Bubble", and "Jawbusters".

(Q14) What are the top 5 all time favorite candy types out of this set?

To solve this in a similar manner, we can use the `tail()` function.

```
# Let's experiment with the dplyr package!
library(dplyr)
candy %>% arrange(winpercent) %>% tail(5)
```

	chocolate	fruity	caramel	peanut	almond	nougat
Snickers	1	0	1		1	1
Kit Kat	1	0	0		0	0
Twix	1	0	1		0	0
Reese's Miniatures	1	0	0		1	0
Reese's Peanut Butter cup	1	0	0		1	0

	crisped	rice	wafer	hard	bar	pluribus	sugar	percent
Snickers				0	0	1		0.546
Kit Kat				1	0	1		0.313
Twix				1	0	1		0.546
Reese's Miniatures				0	0	0		0.034
Reese's Peanut Butter cup				0	0	0		0.720

	price	percent	winpercent
Snickers	0.651		76.67378
Kit Kat	0.511		76.76860

According to the data, the 5 most liked candies are "Reese's Peanut Butter Cups", "Reese's Miniatures", "Twix", "Kit Kat", and "Snickers".

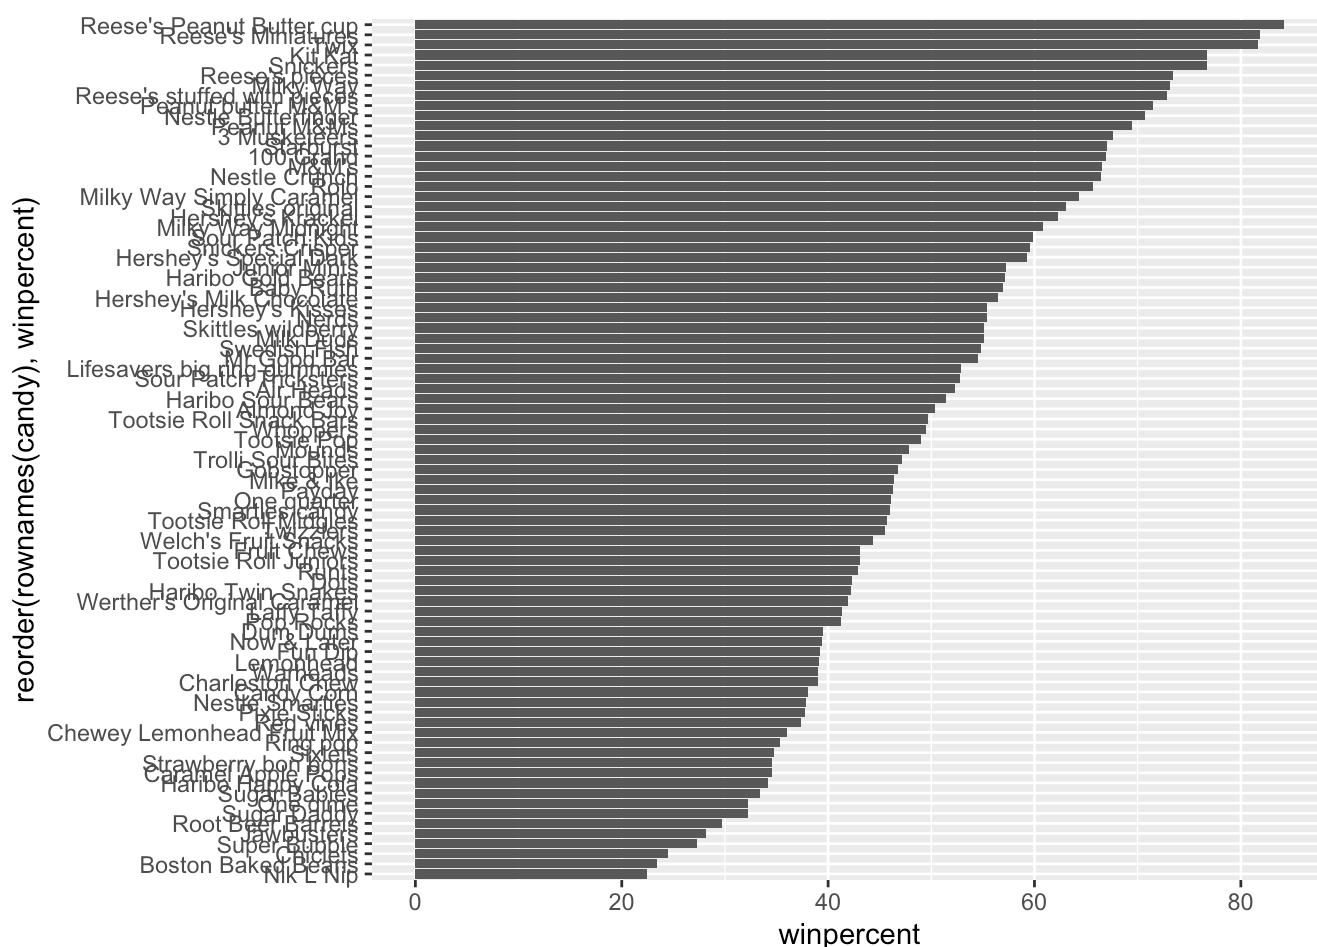
Let's make an intricate barplot to analyze the overall rankings of the candy. We will start simple and build from there.

```
library(ggplot2)
ggplot(candy) +
  aes(winpercent, rownames(candy)) +
  geom_col()
```



7/14

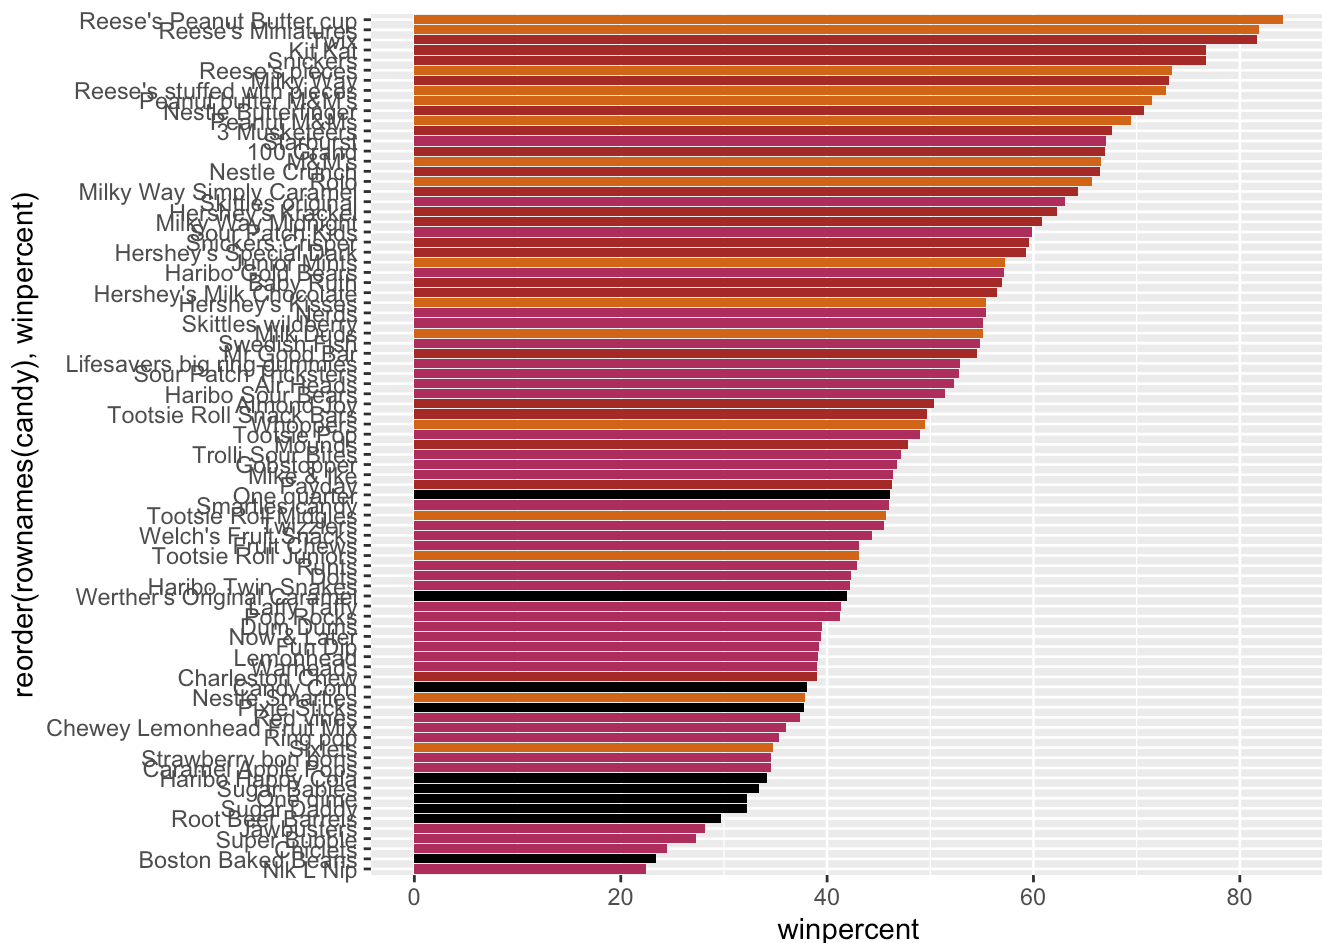
```
ggplot(candy) +
  aes(winpercent, reorder(rownames(candy), winpercent)) +
  geom_col()
```



I can make it even prettier by adding colors to the bars. Let's color by candy type (*the code is borrowed from the Class Lab worksheet*), and reprint the barplot.

```
# Here we define our color categories,
my_cols=rep("black", nrow(candy))
my_cols[as.logical(candy$chocolate)] = "chocolate"
my_cols[as.logical(candy$bar)] = "brown"
my_cols[as.logical(candy$fruity)] = "maroon"

# And here we reprint our barplot.
ggplot(candy) +
  aes(winpercent, reorder(rownames(candy),winpercent)) +
  geom_col(fill=my_cols)
```

```
# Why not save it?
ggsave("my_ugly_barplot.png", width=6, height=9)
```

Using our new colored plot...

(Q17) What is the worst ranked chocolate candy?

The lowest ranked chocolate candy is "Sixlets".

(Q18) What is the best ranked fruity candy?

The highest ranked fruity candy is "Starburst".

Taking a look at pricepercent

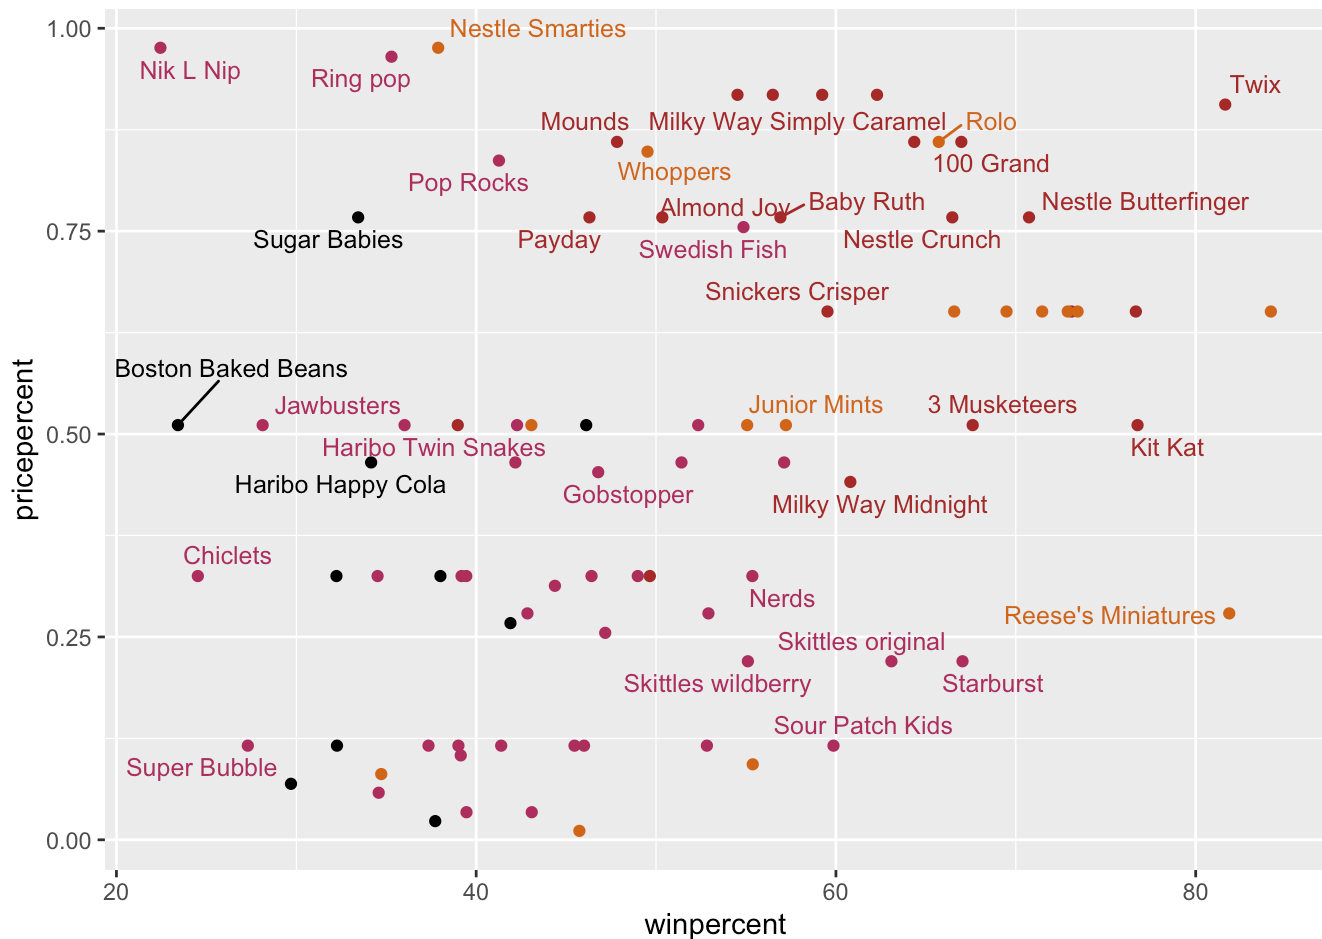
Let's compare pricepercent and winpercent to see which candy will give us the greatest bang for our buck. We want lower pricepercent and higher winpercent.

```
library(ggrepel)

# How about a plot of price vs win?
```

```
ggplot(candy) +
  aes(winpercent, pricepercent, label=rownames(candy)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols, size=3.3, max.overlaps = 5)
```

Warning: ggrepel: 50 unlabeled data points (too many overlaps). Consider increasing max.overlaps



(Q19) Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

```
ord <- order(candy$winpercent, decreasing = TRUE)
head(candy[ord,c(11,12)], n=5)
```

	pricepercent	winpercent
Reese's Peanut Butter cup	0.651	84.18029
Reese's Miniatures	0.279	81.86626
Twix	0.906	81.64291
Kit Kat	0.511	76.76860
Snickers	0.651	76.67378

We can see that the best candies with regards to both **pricepercent** and **winpercent** are **Reese's Miniatures**.

(Q20) What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

```
ord <- order(candy$pricepercent, decreasing = TRUE)
head(candy[ord,c(11,12)], n=5)
```

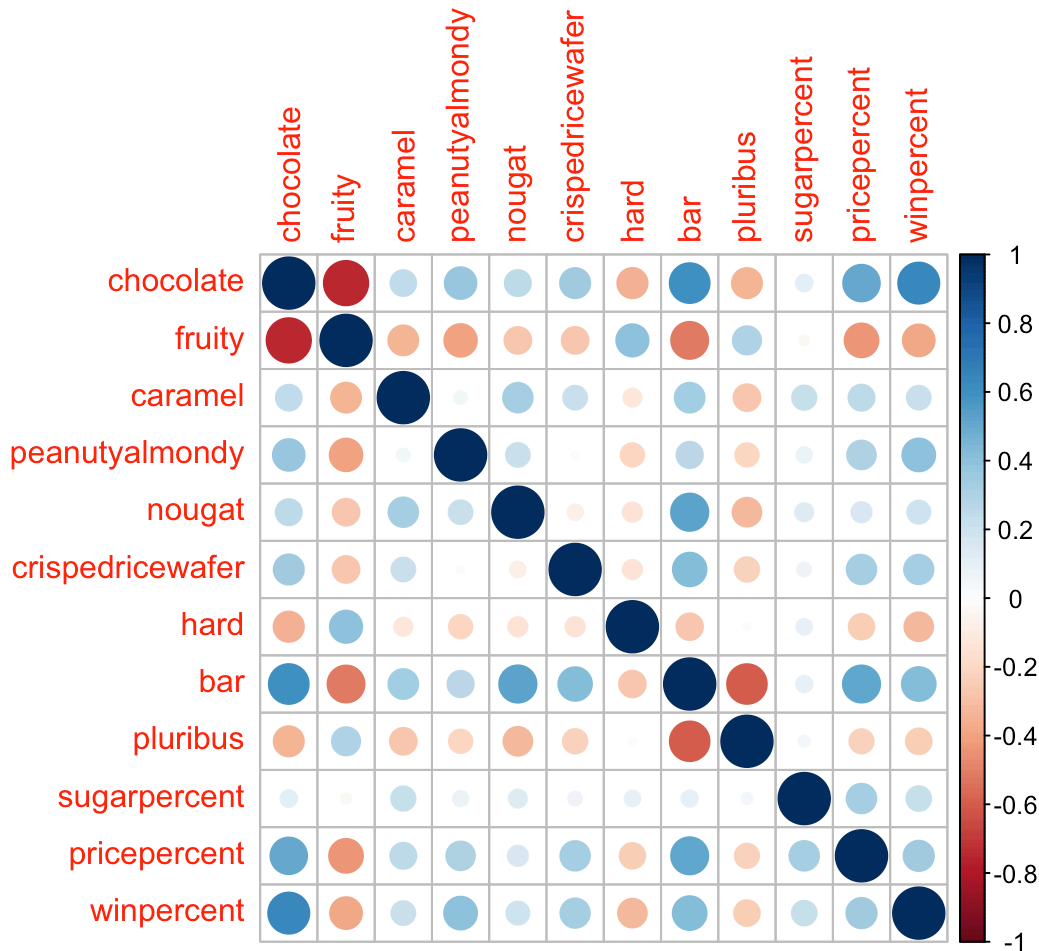
	pricepercent	winpercent
Nik L Nip	0.976	22.44534
Nestle Smarties	0.976	37.88719
Ring pop	0.965	35.29076
Hershey's Krackel	0.918	62.28448
Hershey's Milk Chocolate	0.918	56.49050

The top 5 most expensive candy types are **Nik L Nip, Nestle Smarties, Ring Pops, Hershey's Krackel, and Hershey's Milk Chocolate**. Out of these 5, the least popular is **Nik L Nip**, which is also the most expensive! Not a great choice...

Exploring the Correlation Structure

Let's check the correlation between different candy characteristics to see what we can find.

```
# We have to download the `corrplot` library first
c <- cor(candy)
corrplot::corrplot(c)
```



(Q22) Examining this plot what two variables are anti-correlated (i.e. have minus values)?

The strongest anti-correlation on this plot is between **chocolate** and **fruity**, with a **-0.74** correlation.

(Q23) Similarly, what two variables are most positively correlated?

Outside of the 1.00 correlation between the same variable, the greatest positive correlation is between **chocolate** and **winpercent**, with a **+0.64** correlation.

Principal Component Analysis

We need to be sure to scale our data, because we recall that **winpercent** is on a different scale than the rest of the data.

```
# Our first two PCAs only account for 47%...
pca <- prcomp(candy, scale=T)
summary(pca)
```

Importance of components:

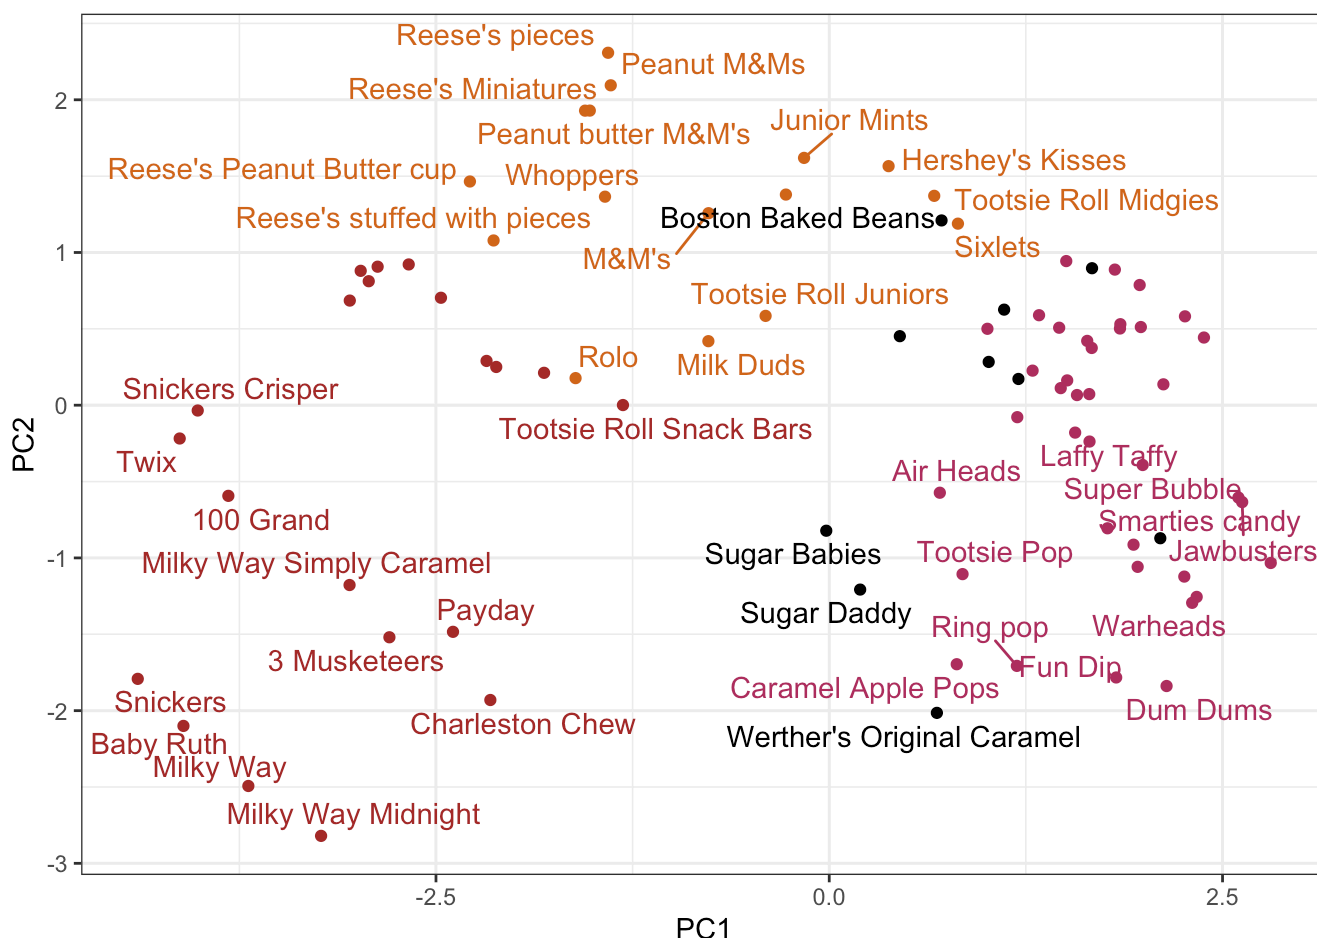
	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530

Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369
	PC8	PC9	PC10	PC11	PC12		
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760		
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317		
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000		

First, lets make a "PCA plot":

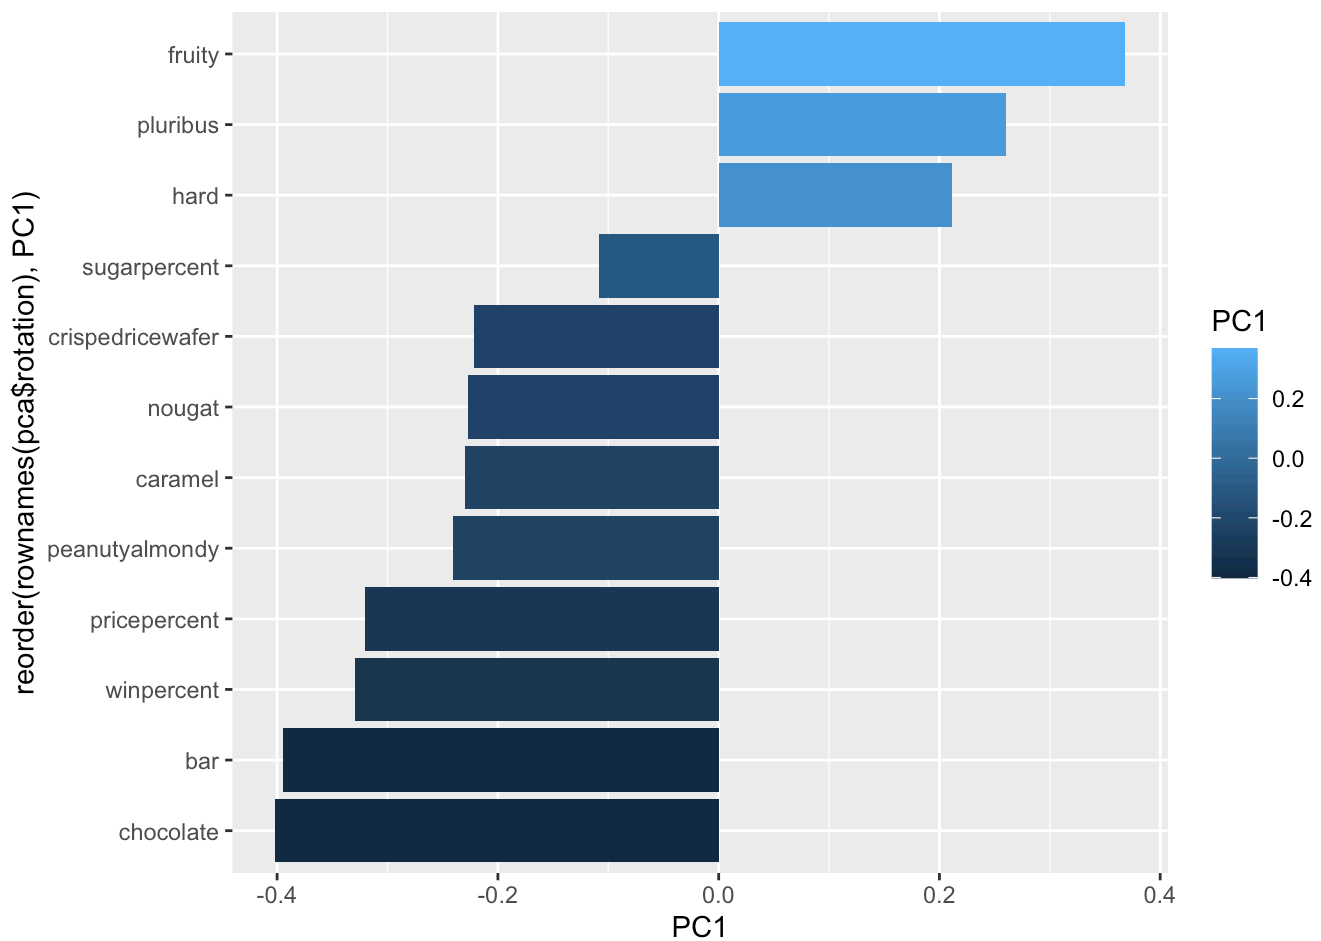
```
ggplot(pca$x) +
  aes(PC1, PC2, label=rownames(pca$x)) +
  geom_point(col=my_cols) +
  geom_text_repel(max.overlaps = 6, col=my_cols) +
  theme_bw()
```

Warning: ggrepel: 43 unlabeled data points (too many overlaps). Consider increasing max.overlaps



The second main PCA result is in the `pca$rotation` object; we can plot this to generate a so-called "loadings" plot.

```
ggplot(pca$rotation) +
  aes(PC1, reorder(rownames(pca$rotation), PC1), fill=PC1) +
  geom_col()
```



(Q24) What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you?

The strongest positively correlated variables with PC1 are **fruity, pluribus, and hard**. This makes sense, because many fruity candies come in packs with more than one candy, and are often also hard candies.