

# Class 13 | RNASeq Analysis

AUTHOR

Christopher Leone | A16731724

## Background

---

Today we will analyze some RNA sequencing data on the effects of a common steroid drug on airway cell lines.

There are two main inputs we need for this analysis:

- **countData**: counts for genes in rows with experiments in the columns
- **colData**: the metadata that tells us about the design of the experiment.

```
# Let's (1) load the libraries:
library(BiocManager)
library(DESeq2)

# And (2) import the files:
counts <- read.csv("airway_scaledcounts.csv", row.names=1)
metadata <- read.csv("airway_metadata.csv")
```

(Q1): How many genes are in this dataset?

```
nrow(counts)
```

```
[1] 38694
```

There are **38,694 genes** in this dataset.

(Q2): How many control cell lines do we have?

```
table(metadata$dex)
```

```
control treated
      4       4
```

There are 4 control cell lines in this table.

## Toy Differential Gene Expression

---

Let's try finding the average or mean of the "control" and "treated" columns and see if they differ.

- First, we need to find all "control" columns
- We need to extract just those columns
- Calculate the `mean()` for each gene "control" values

```
# I like the dplyr system, so I will use it here:
library(dplyr)
```

```
# (1) Filtering for "control" rows only:
control <- metadata %>% filter(dex=="control")
control.counts <- counts %>% select(control$id)

# (2) Taking and storing the means, displaying the head
control.mean <- rowSums(control.counts)/4
head(control.mean)
```

```
ENSG00000000003 ENSG00000000005 ENSG00000000419 ENSG00000000457 ENSG00000000460
          900.75           0.00           520.50           339.75           97.25
ENSG000000000938
          0.75
```

(Q3): Do the same for "treated: to get a `treated.mean`.

```
# (1) Filtering for "treated" rows only:
treated <- metadata %>% filter(dex=="treated")
treated.counts <- counts %>% select(treated$id)

# (2) Taking and storing the means, displaying the head
treated.mean <- rowSums(treated.counts)/4
head(treated.mean)
```

```
ENSG00000000003 ENSG00000000005 ENSG00000000419 ENSG00000000457 ENSG00000000460
          658.00           0.00           546.00           316.50           78.75
ENSG000000000938
          0.00
```

(Q4): And create a plot of `control.mean` vs `treated.mean`.

Ultimately, I decided to put this on logarithmic axes due to the original plot showing most points overlapping. This gives us a much more useful plot.


```
# Let's load the library and make a DF:
library(ggplot2)
means <- data.frame(control.mean, treated.mean)

# Then make the plot:
ggplot(means) +
  aes(x=control.mean, y=treated.mean) +
```

```
geom_point(alpha=0.5) +
scale_x_log10() +
scale_y_log10()
```

Warning in scale\_x\_log10(): log-10 transformation introduced infinite values.

Warning in scale\_y\_log10(): log-10 transformation introduced infinite values.



A common “rule-of-thumb” is to focus on genes with a log2 “fold-change” of  $\pm 2$ . This would indicate a significant up/down regulation.

What if we wanted our axes on a `log2()` scale? Let’s change our previous plot:

```
# We will add a log2 fold-change to a table and make a base plot.
means$log2 <- log2(means$treated.mean/means$control.mean)
```

(Q5): Remove any “zero count” genes from our dataset for further analysis

We end up with 21817 genes that do not have any zero values.

```
# We have to omit anything with zero values
to.keep <- rowSums(means[,1:2] == 0) == 0
```

```
mycounts <- means[to.keep,]
head(mycounts)
```

	control.mean	treated.mean	log2
ENSG00000000003	900.75	658.00	-0.45303916
ENSG000000000419	520.50	546.00	0.06900279
ENSG000000000457	339.75	316.50	-0.10226805
ENSG000000000460	97.25	78.75	-0.30441833
ENSG000000000971	5219.00	6687.50	0.35769358
ENSG000000001036	2327.00	1785.75	-0.38194109

(Q6): How many genes are upregulated? What about downregulated?

There are 314 upregulated genes (according to our parameters), and 485 downregulated genes.

```
sum(mycounts$log2 >= 2)
```

```
[1] 314
```

```
sum(mycounts$log2 <= -2)
```

```
[1] 485
```

## DESeq2 Analysis

Let's do this properly and consider the stats—are the differences in the means significant? Let's use **DESeq2** for this.

The first function we will use from this package sets up the input in the particular format that DESeq wants.

```
dds <- DESeqDataSetFromMatrix(countData = counts,
                              colData = metadata,
                              design = ~dex)
```

converting counts to integer mode

Warning in DESeqDataSet(se, design = design, ignoreRank): some variables in design formula are characters, converting to factors

We can now run our DESeq analysis:

```
dds <- DESeq(dds)
```

estimating size factors

estimating dispersions

gene-wise dispersion estimates

mean-dispersion relationship

final dispersion estimates

fitting model and testing

```
res <-results(dds)
head(res)
```

log2 fold change (MLE): dex treated vs control

Wald test p-value: dex treated vs control

DataFrame with 6 rows and 6 columns

	baseMean	log2FoldChange	lfcSE	stat	pvalue
	<numeric>	<numeric>	<numeric>	<numeric>	<numeric>
ENSG00000000003	747.194195	-0.3507030	0.168246	-2.084470	0.0371175
ENSG00000000005	0.000000	NA	NA	NA	NA
ENSG000000000419	520.134160	0.2061078	0.101059	2.039475	0.0414026
ENSG000000000457	322.664844	0.0245269	0.145145	0.168982	0.8658106
ENSG000000000460	87.682625	-0.1471420	0.257007	-0.572521	0.5669691
ENSG000000000938	0.319167	-1.7322890	3.493601	-0.495846	0.6200029
	padj				
	<numeric>				
ENSG00000000003	0.163035				
ENSG00000000005	NA				
ENSG000000000419	0.176032				
ENSG000000000457	0.961694				
ENSG000000000460	0.815849				
ENSG000000000938	NA				

## Result Figure: Volcano Plots


Here, we will be plotting the adjusted P-values (**padj**) vs the **log2fc**. We are looking for very small P-values.

```
# To color the points:
mycols = rep("black", nrow(res))
mycols[res$log2FoldChange <= -2] <- "blue"
mycols[res$log2FoldChange >= 2] <- "blue"
mycols[res$padj >= 0.05] <- "black"

# The plot:
plot(res$log2FoldChange, -log(res$padj), col=mycols)

# Finally, the boundaries of significant stats:
abline(v=2, col="red")
```

```
abline(v=-2, col="red")
abline(h=-log(0.05), col="red")
```




Let's do better with a `ggplot()` :

```
ggplot(as.data.frame(res)) +
  aes(x=log2FoldChange, y=-log(padj)) +
  geom_point(alpha=0.4, col=mycols) +
  geom_vline(xintercept = c(-2, 2), col="darkgray") +
  geom_hline(yintercept = -log(0.05), col="darkgray") +
  theme_bw() +
  labs(title = "Volcano Plot | DESeq Analysis Results",
       x="Log2 Fold Change",
       y="Adjusted P-value")
```

Warning: Removed 23549 rows containing missing values or values outside the scale range (`geom\_point()`).

## Volcano Plot | DESeq Analysis Results



## Continued: Class 14

What if we want to add gene symbols so we know what genes we are dealing with? We first need to "translate" between the ENSEMBL identifiers to their respective genes.

Let's load some new packages from [BiocManager](#):

```
library(AnnotationDbi)
```

Attaching package: 'AnnotationDbi'

The following object is masked from 'package:dplyr':

select

```
library(org.Hs.eg.db)
```

What different database ID types can I translate between?

```
# We will want to translate between ENSEMBL and SYMBOL.
columns(org.Hs.eg.db)
```

```
[1] "ACCNUM"      "ALIAS"      "ENSEMBL"    "ENSEMBLPROT" "ENSEMBLTRANS"
[6] "ENTREZID"    "ENZYME"     "EVIDENCE"   "EVIDENCEALL"  "GENENAME"
[11] "GENETYPE"    "GO"         "GOALL"      "IPI"           "MAP"
[16] "OMIM"        "ONTOLOGY"   "ONTOLOGYALL" "PATH"          "PFAM"
[21] "PMID"        "PROSITE"    "REFSEQ"     "SYMBOL"        "UCSCKG"
[26] "UNIPROT"
```

So, we need to map between ENSEMBL and SYMBOL ID types to get the data we want. We will also add two more translations to show the gene names and their ENTREZ IDs.

```
# Translation to gene symbols:
res$symbol <- mapIds(x=org.Hs.eg.db,
                    keys=rownames(res),
                    keytype="ENSEMBL",
                    column="SYMBOL")
```

'select()' returned 1:many mapping between keys and columns

```
# Translation to gene names:
res$genename <- mapIds(x=org.Hs.eg.db,
                      keys=rownames(res),
                      keytype="ENSEMBL",
                      column="GENENAME")
```

'select()' returned 1:many mapping between keys and columns

```
# Translation to gene ENTREZ IDs:
res$entrezid <- mapIds(x=org.Hs.eg.db,
                     keys=rownames(res),
                     keytype="ENSEMBL",
                     column="ENTREZID")
```

'select()' returned 1:many mapping between keys and columns

```
# And a preview:
head(res)
```

log2 fold change (MLE): dex treated vs control

Wald test p-value: dex treated vs control

DataFrame with 6 rows and 9 columns

	baseMean	log2FoldChange	lfcSE	stat	pvalue
	<numeric>	<numeric>	<numeric>	<numeric>	<numeric>
ENSG000000000003	747.194195	-0.3507030	0.168246	-2.084470	0.0371175
ENSG000000000005	0.000000	NA	NA	NA	NA
ENSG0000000000419	520.134160	0.2061078	0.101059	2.039475	0.0414026



ENSG00000000457	322.664844	0.0245269	0.145145	0.168982	0.8658106
ENSG00000000460	87.682625	-0.1471420	0.257007	-0.572521	0.5669691
ENSG00000000938	0.319167	-1.7322890	3.493601	-0.495846	0.6200029
	padj	symbol	genename	entrezid	
	<numeric>	<character>	<character>	<character>	
ENSG00000000003	0.163035	TSPAN6	tetraspanin 6	7105	
ENSG00000000005	NA	TNMD	tenomodulin	64102	
ENSG00000000419	0.176032	DPM1	dolichyl-phosphate m..	8813	
ENSG00000000457	0.961694	SCYL3	SCY1 like pseudokina..	57147	
ENSG00000000460	0.815849	FIRRM	FIGNL1 interacting r..	55732	
ENSG00000000938	NA	FGR	FGR proto-oncogene, ..	2268	

Be sure to save our annotated results to a file!

```
write.csv(res, file="my_annotated_results.csv")
```

## Pathway Analysis

First, we need to load some new packages from our library.

```
library(pathview)
```

#####  
 Pathview is an open source software package distributed under GNU General Public License version 3 (GPLv3). Details of GPLv3 is available at <http://www.gnu.org/licenses/gpl-3.0.html>. Particullary, users are required to formally cite the original Pathview paper (not just mention it) in publications or products. For details, do citation("pathview") within R.

The pathview downloads and uses KEGG data. Non-academic uses may require a KEGG license agreement (details at <http://www.kegg.jp/kegg/legal.html>).

#####

```
library(gage)
```

```
library(gageData)
```

Let's examine the gageData:

```
data(kegg.sets.hs)
```

```
# Examine the first 2 pathways in this kegg set for humans
head(kegg.sets.hs, 2)
```

```
$`hsa00232 Caffeine metabolism`
[1] "10" "1544" "1548" "1549" "1553" "7498" "9"
```

```
$`hsa00983 Drug metabolism - other enzymes`
```

```
[1] "10"      "1066"    "10720"   "10941"   "151531"  "1548"    "1549"    "1551"
[9] "1553"    "1576"    "1577"    "1806"    "1807"    "1890"    "221223"  "2990"
[17] "3251"    "3614"    "3615"    "3704"    "51733"   "54490"   "54575"   "54576"
[25] "54577"   "54578"   "54579"   "54600"   "54657"   "54658"   "54659"   "54963"
[33] "574537"  "64816"   "7083"    "7084"    "7172"    "7363"    "7364"    "7365"
[41] "7366"    "7367"    "7371"    "7372"    "7378"    "7498"    "79799"   "83549"
[49] "8824"    "8833"    "9"       "978"
```

To run pathway analysis, we will use the `gage()` function and it requires a "vector of importance". We will use our `log2FC` results from `res`.

```
foldchanges = res$log2FoldChange
names(foldchanges) = res$entrezid
head(foldchanges)
```

```
      7105      64102      8813      57147      55732      2268
-0.35070302      NA  0.20610777  0.02452695 -0.14714205 -1.73228897
```

```
# And finally, let's run gage!
keggres = gage(foldchanges, gsets=kegg.sets.hs)
```

What is in these results?

```
attributes(keggres)
```

```
$names
[1] "greater" "less"    "stats"
```

```
# What pathways overlap with what we have annotated?
head(keggres$less)
```

	p.geomean	stat.mean
hsa05332 Graft-versus-host disease	0.0004250461	-3.473346
hsa04940 Type I diabetes mellitus	0.0017820293	-3.002352
hsa05310 Asthma	0.0020045888	-3.009050
hsa04672 Intestinal immune network for IgA production	0.0060434515	-2.560547
hsa05330 Allograft rejection	0.0073678825	-2.501419
hsa04340 Hedgehog signaling pathway	0.0133239547	-2.248547
	p.val	q.val
hsa05332 Graft-versus-host disease	0.0004250461	0.09053483
hsa04940 Type I diabetes mellitus	0.0017820293	0.14232581
hsa05310 Asthma	0.0020045888	0.14232581
hsa04672 Intestinal immune network for IgA production	0.0060434515	0.31387180
hsa05330 Allograft rejection	0.0073678825	0.31387180
hsa04340 Hedgehog signaling pathway	0.0133239547	0.47300039
	set.size	exp1
hsa05332 Graft-versus-host disease	40	0.0004250461

hsa04940 Type I diabetes mellitus	42	0.0017820293
hsa05310 Asthma	29	0.0020045888
hsa04672 Intestinal immune network for IgA production	47	0.0060434515
hsa05330 Allograft rejection	36	0.0073678825
hsa04340 Hedgehog signaling pathway	56	0.0133239547

And let's use the `pathview()` function to look at what the pathway looks like! Let's look at asthma, since that was what we have been investigating.

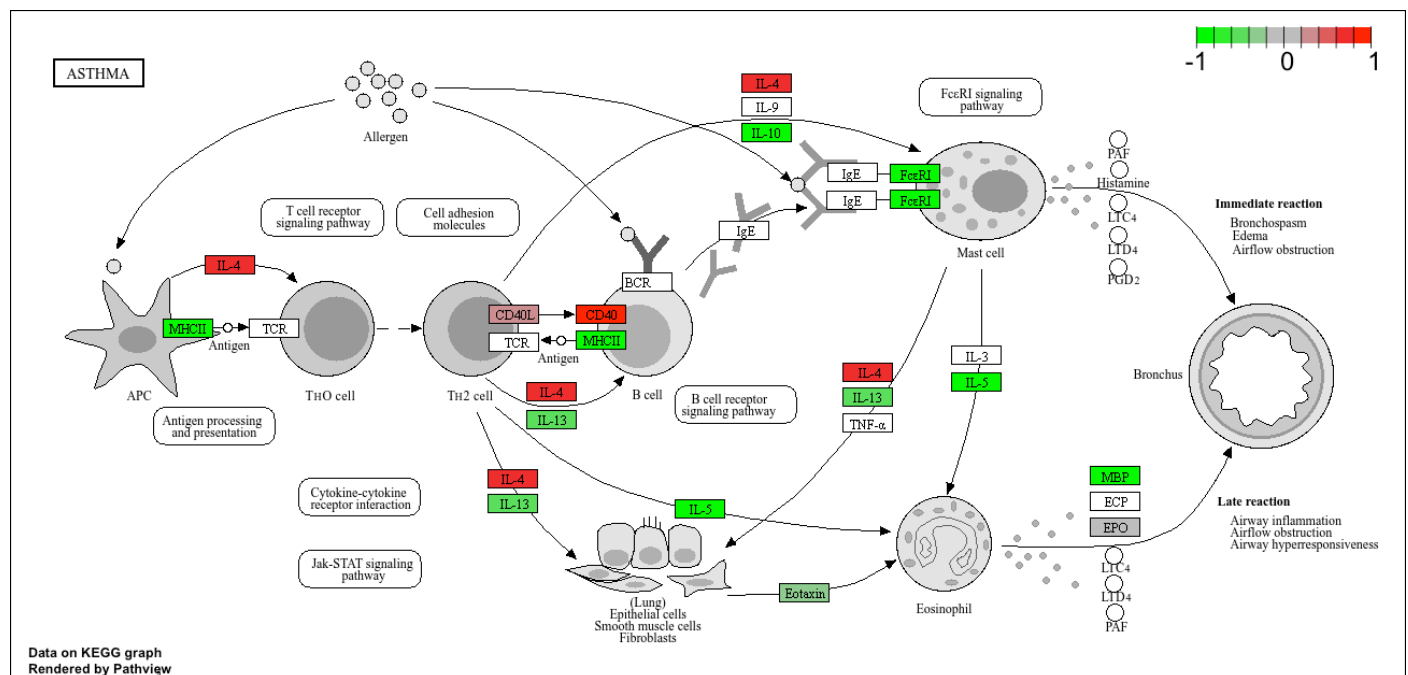
```
pathview(gene.data=foldchanges, pathway.id="hsa05310")
```

'select()' returned 1:1 mapping between keys and columns

Info: Working in directory /Users/chrisleone/Desktop/BIMM 143 | Rstudio/class13

Info: Writing image file hsa05310.pathview.png

The image that was generated can be seen here:



The Asthma pathway overlaps with our differentially expressed genes.